



Information Retrieval and Web Search

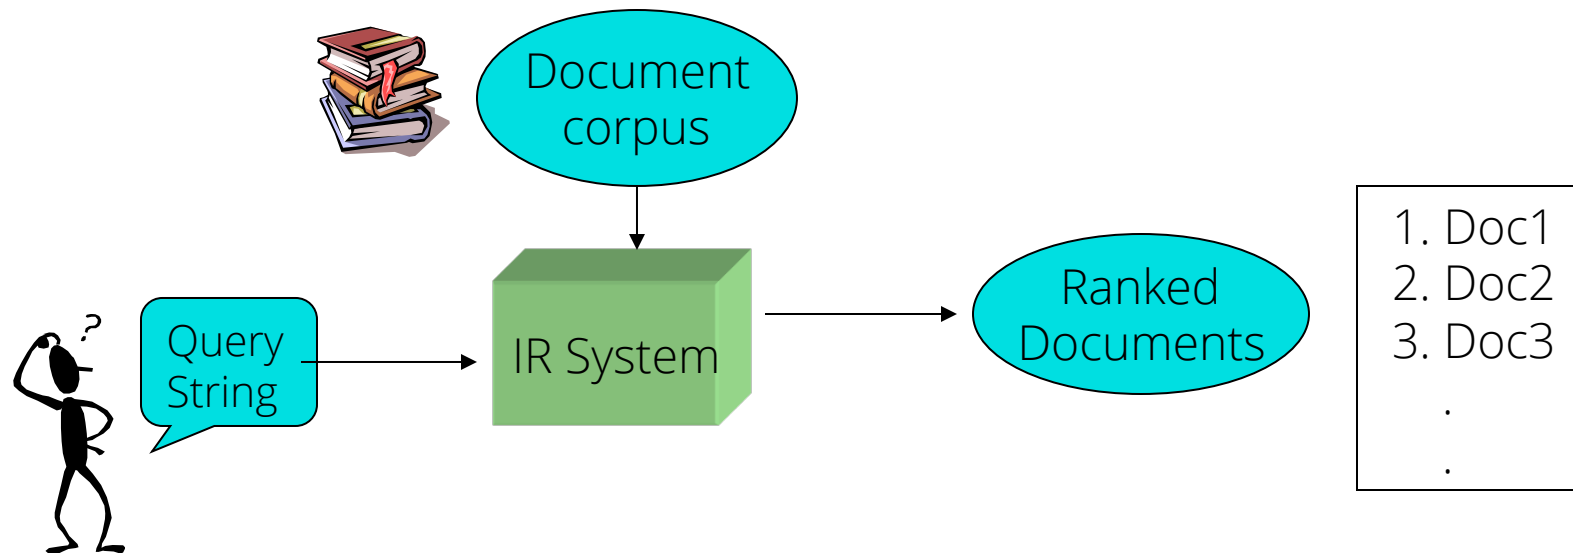
Boolean retrieval

Instructor: Rada Mihalcea

(Note: some of the slides in this set have been adapted from a course taught by Prof. Chris Manning at Stanford University)

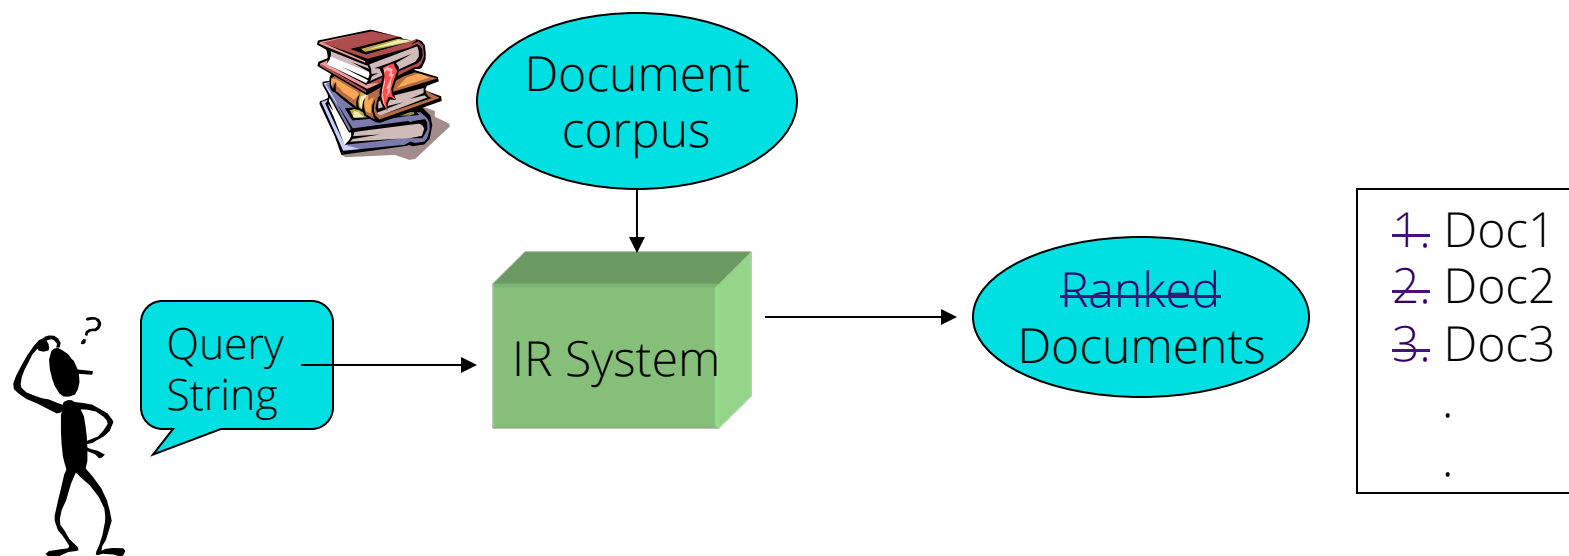
Typical IR task

- Input:
 - A large collection of unstructured text documents.
 - A user query expressed as text.
- Output:
 - A ranked list of documents that are relevant to the query.



Boolean ~~Typical~~ IR task

- Input:
 - A large collection of unstructured text documents.
 - A user query expressed as text.
- Output:
 - A ~~ranked~~ list of documents that are relevant to the query.



Boolean retrieval

- Information Need: *Which plays by Shakespeare mention Brutus and Caesar, but not Calpurnia?*
- Boolean Query: Brutus AND Caesar AND NOT Calpurnia
- Possible search procedure:
 - Linear scan through all documents (Shakespeare's collected works).
 - Compile list of documents that contain Brutus and Caesar, but not Calpurnia.
 - Advantage: simple, it works for moderately sized corpora.
 - Disadvantage: need to do linear scan for every query \Rightarrow slow for large corpora.

Term-document incidence matrices

- Precompute a data structure that makes search fast for every query.

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0

1 if document contains word, 0 otherwise

Term-document incidence matrix M

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0

Query = Brutus AND Caesar AND NOT Calpurnia

Answer = $M(\text{Brutus}) \wedge M(\text{Caesar}) \wedge \neg M(\text{Calpurnia})$
 $= 1\ 1\ 0\ 1\ 0\ 0 \wedge 1\ 1\ 0\ 1\ 1\ 1 \wedge 1\ 0\ 1\ 1\ 1\ 1$
 $= 1\ 0\ 0\ 1\ 0\ 0$
 \Rightarrow Anthony and Cleopatra, Hamlet

110100 \wedge
 110111 \wedge
 101111
 100100

Answers to Query

- **Antony and Cleopatra**, Act III, Scene ii

Agrippa [Aside to DOMITIUS ENOBARBUS]: Why, Enobarbus,





When Antony found Julius *Caesar* dead,
He cried almost to roaring; and he wept
When at Philippi he found *Brutus* slain.

- **Hamlet**, Act III, Scene ii

Lord Polonius: I did enact Julius *Caesar* I was killed i' the
Capitol; *Brutus* killed me.

Scalability: Dense Format

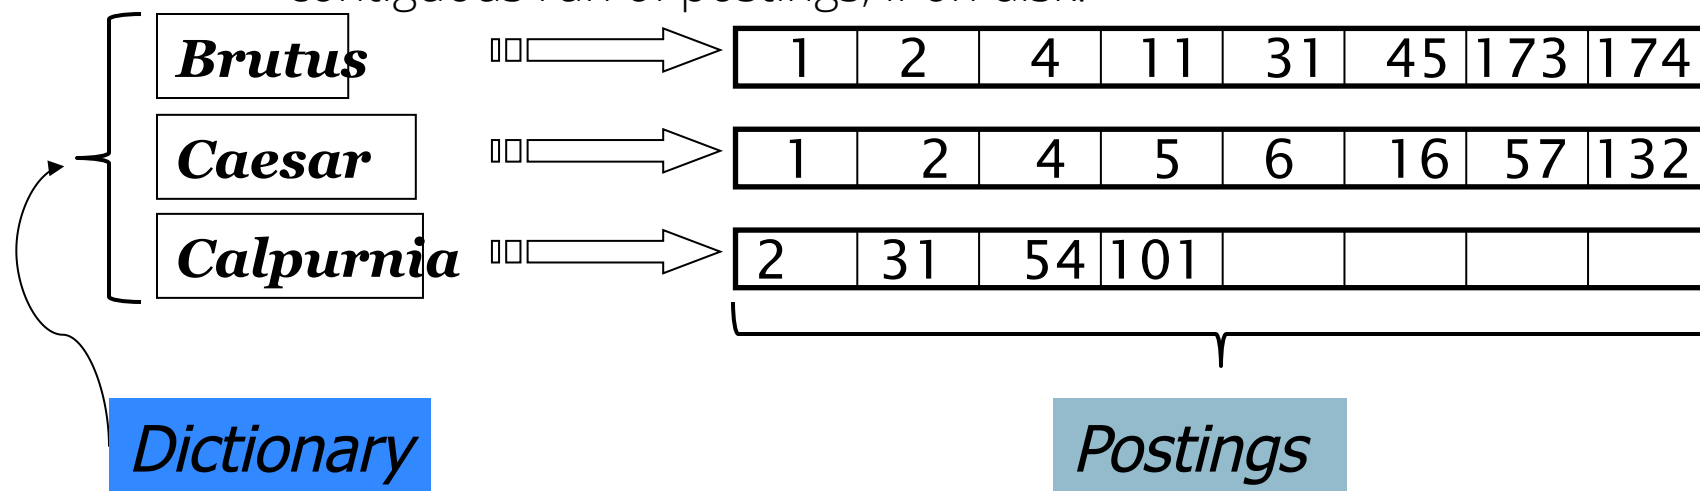
- Assume:
 - Corpus has 1 million documents.
 - Each document is about 1,000 words long.
 - Each word takes 6 bytes, on average.
 - Of the 1 billion word tokens 500,000 are unique.
- Then:
 - Corpus storage takes:
 - $1M * 1,000 * 6 = 6GB$ 
 - Term-Document incidence matrix would take:
 - $500,000 * 1,000,000 = 0.5 * 10^{12}$ bits 

Scalability: Sparse Format

- Of the 500 billion entries, at most 1 billion are non-zero.
 - ⇒ at least 99.8% of the entries are zero.
 - ⇒ use a sparse representation to reduce storage size!
- Store only non-zero entries ⇒ **Inverted Index**.

Inverted Index for Boolean Retrieval

- Map each term to a **posting list** of documents containing it:
 - Identify each document by a numerical **docID**.
 - **Dictionary** of **terms** usually in memory.
 - Posting list:
 - linked lists of variable-sized array, if in memory.
 - contiguous run of postings, if on disk.



Inverted Index: Step 1

- Assemble sequence of $\langle \text{token}, \text{docID} \rangle$ pairs.
 - assume text has been tokenized

Doc 1

I did enact Julius
Caesar I was killed
i' the Capitol;
Brutus killed me.

Doc 2

So let it be with
Caesar. The noble
Brutus hath told you
Caesar was ambitious



Term	docID
I	1
did	1
enact	1
julius	1
caesar	1
I	1
was	1
killed	1
i'	1
the	1
capitol	1
brutus	1
killed	1
me	1
so	2
let	2
it	2
be	2
with	2
caesar	2
the	2
noble	2
brutus	2
hath	2
told	2
you	2
caesar	2
was	2
ambitious	2

Inverted Index: Step 2

- Sort by terms, then by docIDs.

Term	docID
I	1
did	1
enact	1
julius	1
caesar	1
I	1
was	1
killed	1
i'	1
the	1
capitol	1
brutus	1
killed	1
me	1
so	2
let	2
it	2
be	2
with	2
caesar	2
the	2
noble	2
brutus	2
hath	2
told	2
you	2
caesar	2
was	2
ambitious	2



Term	docID
ambitious	2
be	2
brutus	1
brutus	2
capitol	1
caesar	1
caesar	2
caesar	2
did	1
enact	1
hath	1
I	1
I	1
i'	1
it	2
julius	1
killed	1
killed	1
let	2
me	1
noble	2
so	2
the	1
the	2
told	2
you	2
was	1
was	2
with	2

Inverted Index: Step 3

- Merge multiple term entries per document.
- Split into **dictionary** and **posting lists**.
 - keep posting lists sorted, for efficient query processing.
- Add **document frequency** information:
 - useful for efficient query processing.
 - also useful later in document ranking.

Inverted Index: Step 3


Term	docID
ambitious	2
be	2
brutus	1
brutus	2
capitol	1
caesar	1
caesar	2
caesar	2
did	1
enact	1
hath	1
I	1
I	1
i'	1
it	2
julius	1
killed	1
killed	1
let	2
me	1
noble	2
so	2
the	1
the	2
told	2
you	2
was	1
was	2
with	2



term	doc. freq.	→	postings lists
ambitious	1	→	2
be	1	→	2
brutus	2	→	1 → 2
capitol	1	→	1
caesar	2	→	1 → 2
did	1	→	1
enact	1	→	1
hath	1	→	2
i	1	→	1
i'	1	→	1
it	1	→	2
julius	1	→	1
killed	1	→	1
let	1	→	2
me	1	→	1
noble	1	→	2
so	1	→	2
the	2	→	1 → 2
told	1	→	2
you	1	→	2
was	2	→	1 → 2
with	1	→	2

Query Processing: AND

- Consider processing the query:
Brutus AND Caesar
 - Locate *Brutus* in the Dictionary;
 - Retrieve its postings.
 - Locate *Caesar* in the Dictionary;
 - Retrieve its postings.
 - “Merge” the two postings (intersect the document sets):



Query Processing: t1 AND t2

INTERSECT(p_1, p_2)

```
1  answer  $\leftarrow \langle \rangle$ 
2  while  $p_1 \neq \text{NIL}$  and  $p_2 \neq \text{NIL}$ 
3  do if  $\text{docID}(p_1) = \text{docID}(p_2)$ 
4      then  $\text{ADD}(\text{answer}, \text{docID}(p_1))$ 
5           $p_1 \leftarrow \text{next}(p_1)$ 
6           $p_2 \leftarrow \text{next}(p_2)$ 
7      else if  $\text{docID}(p_1) < \text{docID}(p_2)$ 
8          then  $p_1 \leftarrow \text{next}(p_1)$ 
9          else  $p_2 \leftarrow \text{next}(p_2)$ 
10 return answer
```

p_1, p_2 – pointers to posting lists
corresponding to t_1 and t_2
 docID – function that returns the Id
of the document in location
pointed by p_i

Query Processing: t1 OR t2

Union

INTERSECT(p_1, p_2)

```
1  answer  $\leftarrow \langle \rangle$ 
2  while  $p_1 \neq \text{NIL}$  and  $p_2 \neq \text{NIL}$ 
3  do if  $\text{docID}(p_1) = \text{docID}(p_2)$ 
4      then  $\text{ADD}(\text{answer}, \text{docID}(p_1))$ 
5           $p_1 \leftarrow \text{next}(p_1)$ 
6           $p_2 \leftarrow \text{next}(p_2)$ 
7      else if  $\text{docID}(p_1) < \text{docID}(p_2)$ 
8          then  $p_1 \leftarrow \text{next}(p_1)$ 
9          else  $p_2 \leftarrow \text{next}(p_2)$ 
10 return answer
```

p_1, p_2 – pointers to posting lists
corresponding to t_1 and t_2
 docID – function that returns the Id
of the document at position p

$\text{ADD}(\text{answer}, \text{docID}(p_1))$

$\text{ADD}(\text{answer}, \text{docID}(p_2))$

Exercise: Query Processing: NOT

- Exercise: Adapt the pseudocode for the query:

t1 AND NOT t2

e.g., Brutus AND NOT Caesar


- Can we still run through the merge in time $O(\text{length}(p1) + \text{length}(p2))$?

Query Optimization:

What is the best order for query processing?

- Consider a query that is an *AND* of n terms.

Query: *Brutus AND Calpurnia AND Caesar*



- For each of the n terms, get its postings, then *AND* them together.
 - Process in order of increasing freq:
 - start with smallest set, then keep cutting further.
 - use document frequencies stored in the dictionary.
- ⇒ execute the query as (*Calpurnia AND Brutus*) *AND Caesar*

More General Optimization

- E.g., (*madding OR crowd*) AND (*ignoble OR strife*)
- Get frequencies for all terms.
- Estimate the size of each *OR* by the sum of its frequencies (conservative).
- Process in increasing order of *OR* sizes.

Exercise

- Recommend a query processing order for:
 - *(tangerine OR trees) AND (marmalade OR skies) AND (kaleidoscope OR eyes)*
 - which two terms should we process first?

Term	Freq
eyes	213312
kaleidoscope	87009
marmalade	107913
skies	271658
tangerine	46653
trees	316812

Extensions to the Boolean Model

- **Phrase Queries:**

- Want to answer query “Information Retrieval”, as a phrase.
- The concept of phrase queries is one of the few “advanced search” ideas that is easily understood by users.
 - about 10% of web queries are phrase queries.
 - many more are *implicit phrase queries* (e.g. person names).

- **Proximity Queries:**

- Altavista: Python NEAR language
- Google: Python * language
- many search engines use keyword proximity *implicitly*.

Solution 1 for Phrase Queries: Biword Indexes

- Index every two consecutive tokens in the text.
 - Treat each biword (or bigram) as a vocabulary term.
 - The text “*modern information retrieval*” generates **biwords**:
 - *modern information*
 - *information retrieval*
 - **Bigram phrase query** processing is now straightforward.
 - **Longer phrase queries?**
 - Heuristic solution: break them into conjunction of biwords.
 - Query “electrical engineering and computer science”:
 - “electrical engineering” AND “engineering and” AND “and computer” AND “computer science”
 - Without verifying the retrieved docs, can have false positives!

Biword Indexes

- Can have false positives:
 - Unless retrieved docs are verified \Rightarrow increased time complexity.
 - Larger dictionary leads to index blowup:
 - clearly unfeasible for ngrams larger than bigrams.
- \Rightarrow not a standard solution for phrase queries:
- but useful in compound strategies.

Solution 2 for Phrase Queries: Positional Indexes

- In the postings list:
 - for each token *tok*:
 - for each document *docID*:
 - store the positions in which *tok* appears in *docID*.
 - < *be*: 993427;
 - 1*: 7, 18, 33, 72, 86, 231;
 - 2*: 3, 149;
 - 4*: 17, 191, 291, 430, 434;
 - 5*: 363, 367, ... >
 - which documents might contain “to be or not to be”?

Positional Indexes: Query Processing

- Use a merge algorithm at two levels:
 1. Postings level, to find matchings docIDs for query tokens.
 2. Document level, to find consecutive positions for query tokens.
- Extract index entries for each distinct term: *to, be, or, not*.
- Merge their *doc:position* lists to enumerate all positions with “*to be or not to be*”.
 - **to:** 2:1,17,74,222,551; 4:8,16,190,429,433; 7:13,23,191; ...
 - **be:** 1:17,19; 4:17,191,291,430,434; 5:14,19,101; ...
- Same general method for proximity searches.

POSITIONALINTERSECT(p_1, p_2, k)

```
1  answer  $\leftarrow \langle \rangle$ 
2  while  $p_1 \neq \text{NIL}$  and  $p_2 \neq \text{NIL}$ 
3  do if  $\text{docID}(p_1) = \text{docID}(p_2)$ 
4      then  $l \leftarrow \langle \rangle$ 
5           $pp_1 \leftarrow \text{positions}(p_1)$ 
6           $pp_2 \leftarrow \text{positions}(p_2)$ 
7          while  $pp_1 \neq \text{NIL}$ 
8              do while  $pp_2 \neq \text{NIL}$ 
9                  do if  $|\text{pos}(pp_1) - \text{pos}(pp_2)| \leq k$ 
10                     then  $\text{ADD}(l, \text{pos}(pp_2))$ 
11                     else if  $\text{pos}(pp_2) > \text{pos}(pp_1)$ 
12                         then break
13                      $pp_2 \leftarrow \text{next}(pp_2)$ 
14                     while  $l \neq \langle \rangle$  and  $|l[0] - \text{pos}(pp_1)| > k$ 
15                         do  $\text{DELETE}(l[0])$ 
16                     for each  $ps \in l$ 
17                         do  $\text{ADD}(\text{answer}, \langle \text{docID}(p_1), \text{pos}(pp_1), ps \rangle)$ 
18                      $pp_1 \leftarrow \text{next}(pp_1)$ 
19                  $p_1 \leftarrow \text{next}(p_1)$ 
20                  $p_2 \leftarrow \text{next}(p_2)$ 
21             else if  $\text{docID}(p_1) < \text{docID}(p_2)$ 
22                 then  $p_1 \leftarrow \text{next}(p_1)$ 
23                 else  $p_2 \leftarrow \text{next}(p_2)$ 
24 return answer
```

Positional Index: Size

- Need an entry for each occurrence, not just for each document.
- Index size depends on average document size:
 - Average web page has less than 1000 terms.
 - Books, even some poems ... easily 100,000 terms.
 - large documents cause an increase of 2 orders of magnitude.
 - Consider a term with frequency 0.1%:

Document size	Expected postings	Expected entries in positional posting
1000	1	1
100,000	1	100

Positional Index

- A positional index expands postings storage *substantially*.
 - 2 to 4 times as large as a non-positional index
 - compressed, it is between a third and a half of uncompressed raw text.
- Nevertheless, a positional index is now standardly used because of the power and usefulness of phrase and proximity queries:
 - whether used explicitly or implicitly in a ranking retrieval system.

Combined Strategy

- Biword and positional indexes can be fruitfully combined:
 - For particular phrases (*“Michael Jackson”*, *“Britney Spears”*) it is inefficient to keep on merging positional postings lists
 - Even more so for phrases like *“The Who”*. Why?
- 1. Use a phrase index, or a biword index, for certain queries:
 - Queries known to be common based on recent querying behavior.
 - Queries where the individual words are common but the desired phrase is comparatively rare.
- 2. Use a positional index for remaining phrase queries.

Boolean Retrieval vs. Ranked Retrieval

- Professional users prefer Boolean query models:
 - Boolean queries are precise: a document either matches the query or it does not.
 - Greater control and transparency over what is retrieved.
 - Some domains allow an effective ranking criterion:
 - Westlaw returns documents in reverse chronological order.
- Hard to tune precision vs. recall:
 - AND operator tends to produce high precision but low recall.
 - OR operator gives low precision but high recall.
 - Difficult/impossible to find satisfactory middle ground.

Boolean Retrieval vs. Ranked Retrieval

- Need an effective method to rank the matched documents.
 - Give more weight to documents that mention a token several times vs. documents that mention it only once.
 - record **term frequency** in the postings list.
- Web search engines implement ranked retrieval models:
 - Most include at least partial implementations of Boolean models:
 - Boolean operators.
 - Phrase search.
 - Still, improvements are generally focused on free text queries
 - Vector space model